



# Deep Learning

**From Fundamentals to Research-Driven**

**Equipping You with Research Depth and  
Industry Skills**

By:

**Dr. Zohair Ahmed**



 [www.youtube.com/@ZohairAI](https://www.youtube.com/@ZohairAI) 

 [www.begindiscovery.com](http://www.begindiscovery.com)

# Development Environment Setup

---

## 1. Install Python (Choose One Option)

- Option 1: Download & install Python from python.org
- Option 2: Install **Miniconda** or **Anaconda (Full)**
- **Recommendation: Use Manual Python**

## 2. Create a Virtual Environment

- Always create a **venv** for each project
- Keeps dependencies clean and avoids version conflicts
- Activate **venv** → install packages inside it

## 3. Use PIP for Package Management

- pip install package\_name
- Keep packages updated
- Manage requirements through requirements.txt

## 4. Start with Google Colab

- Best for beginners
- Free GPU
- No installation required
- Easier than Jupyter Notebook

## 5. Use Proper IDE for Real Projects

- **VS Code** (recommended)
- or **Spyder**

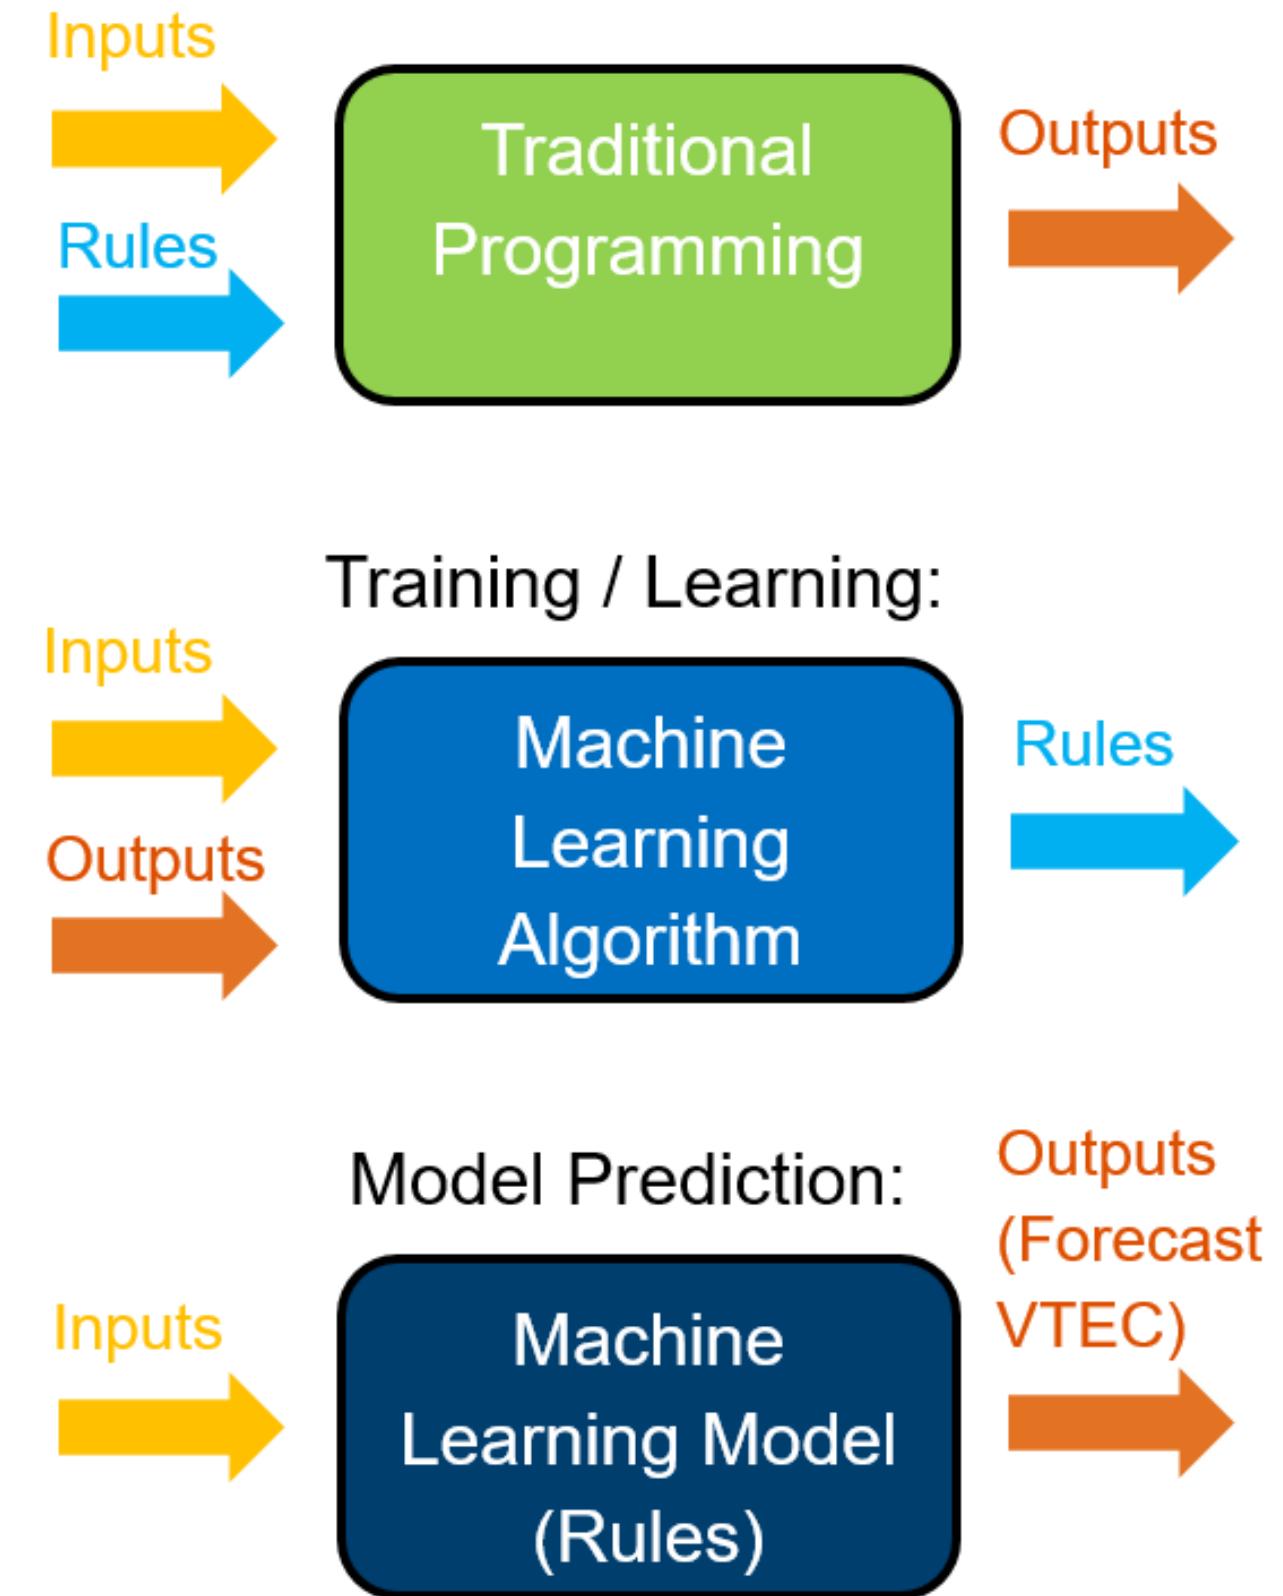
## 6. Build a Kaggle Portfolio

- Join competitions
- Publish notebooks
- Improve skills + visibility
- Essential for **Data Scientist career**



# Why Study Learning Systems?

- Computing has evolved from:
  - *Explicit programming* → hand-crafted rules
  - To *learning from data* → models discover patterns
- Why learning is essential:
  - Many tasks are too complex to hand-code (e.g., vision, language)
  - Data is abundant; rules are not
  - Adaptive systems can improve over time
- Goal: Build systems that improve performance with experience.



# Traditional Programming Example

---

- **Task:** Classify whether a number is even or odd.
- We write rules ourselves:
- **Input:** number
- **Rules:** written by programmer ( $n \% 2 == 0$ )
- **Output:** “Even” or “Odd”

```
def is_even(n):  
    if n % 2 == 0:  
        return "Even"  
    else:  
        return "Odd"
```



# Machine Learning Example

- Task: Predict house prices
- *Here, we do NOT write rules.*
- ML learns rules from data.
- **We provide:**
- **Input:** house size, rooms, location
- **Output:** actual price
- ML finds patterns = “rules”
- The ML algorithm produces something like:
- Price= $w \times \text{Size} + b \rightarrow 0.015 \times \text{Size} - 3$

Size (sqft)	Price
1000	12 lac
1500	18 lac
1800	25 lac



# Everyday Examples of Learning Systems

---

- In your daily life:
- Recommendation systems: YouTube, Netflix, Spotify
- Smart assistants: Siri, Alexa, Google Assistant
- Social media feeds: ranking, content moderation
- Navigation apps: traffic prediction, route planning
- **Key message:** Learning-based methods already mediate much of your digital experience.



# What is Artificial Intelligence (AI)?

---

- **Informal definition:**

- *AI = Systems that perform tasks that, if done by humans, would require intelligence.*

- **Typical AI capabilities:**

- Perception (seeing, hearing)
  - Reasoning and planning
  - Learning and adaptation
  - Natural language understanding and generation

- **Subfields:**

- Knowledge representation, search, planning, robotics, ML, NLP, etc.



# Historical Approaches to AI

---

- **Symbolic / Rule-based AI:**
- Expert systems, logic rules, search-based planning
- Strength: explicit reasoning, interpretability
- Weakness: brittle, hard to scale to messy real-world data
- **Statistical AI:**
- Probability, graphical models, pattern recognition
- Strength: handles uncertainty, data-driven
- Setup for **machine learning** and later **deep learning**.



# What is Machine Learning (ML)?

---

- **Core idea:**
  - ML is about **learning a function** from data, rather than hand-coding it.
  - **Classic definition (Mitchell):**

*A computer program is said to learn from experience  $E$  with respect to some class of tasks  $T$  and performance measure  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .*
  - **Now, the Super Simple Version**
  - Imagine you have a friend named **Robo**.
  - **Task (T):** What you want Robo to do → e.g., recognize cats in pictures.
  - **Experience (E):** The practice Robo gets → e.g., showing Robo many pictures of cats and not-cats.
  - **Performance Measure (P):** How you judge Robo's progress → e.g., how many pictures Robo gets right.
  - **What learning means:**
  - **Robo is learning if:**
    - After seeing more examples (experience  $E$ ),
    - Robo gets better at the task (task  $T$ ),
    - And the score you measure (performance  $P$ ) improves.
  - **Even simpler:**
  - **Learning** = getting better at something by practicing.
  - A program is “**learning**” if it practices something and its results improve.
- 



# Example: Machine Learning (ML)

---

- **Example:**
- **Task (T):** Predicting tomorrow's temperature
- **Experience (E):** Past weather data
- **Performance (P):** How close its predictions are to the real temperature
- If the computer predicts better and better as it sees more data → **It has learned!**
- **Key components:**
  - Data (examples)
  - Model (hypothesis space)
  - Objective / loss function
  - Learning algorithm (optimization)



# What is Deep Learning (DL)?

---

- Deep Learning is a part of Machine Learning where we use **deep neural networks**, meaning **neural networks with many layers**.
- Think of it like teaching a child to recognize things:
  - First they learn **lines**
  - Then **shapes**
  - Then **eyes, nose, ears**
  - Then **faces**
  - Then **who the person is**
- Neural networks do something similar.
- **Why do we call it deep?**
- Because the network has **many layers**, one after another like a tall sandwich.
- **These layers learn hierarchically:**
  - **Low-level layers** → simple things
    - edges
    - lines
    - corners
  - **Mid-level layers** → patterns
    - textures
    - shapes
  - **High-level layers** → big concepts
    - faces
    - cars
    - words
    - meaning in a sentence
  - This “stacking” makes it powerful.



# What is Deep Learning (DL)?

---

- **Machine Learning (ML)**: A broad field: models learn patterns from data to make predictions/decisions.
  - Includes: linear/logistic regression, decision trees, random forests, gradient boosting, SVMs, kNN, Naive Bayes, clustering, etc.
- **Neural Networks (NN)**: One family of ML models made of layers (**linear transform + non-linear activation**).
- **Deep Learning (DL)**: Neural networks with **many layers** (deep architectures), plus training tricks and large-scale data/compute.



# Why Deep Learning is so good?

---

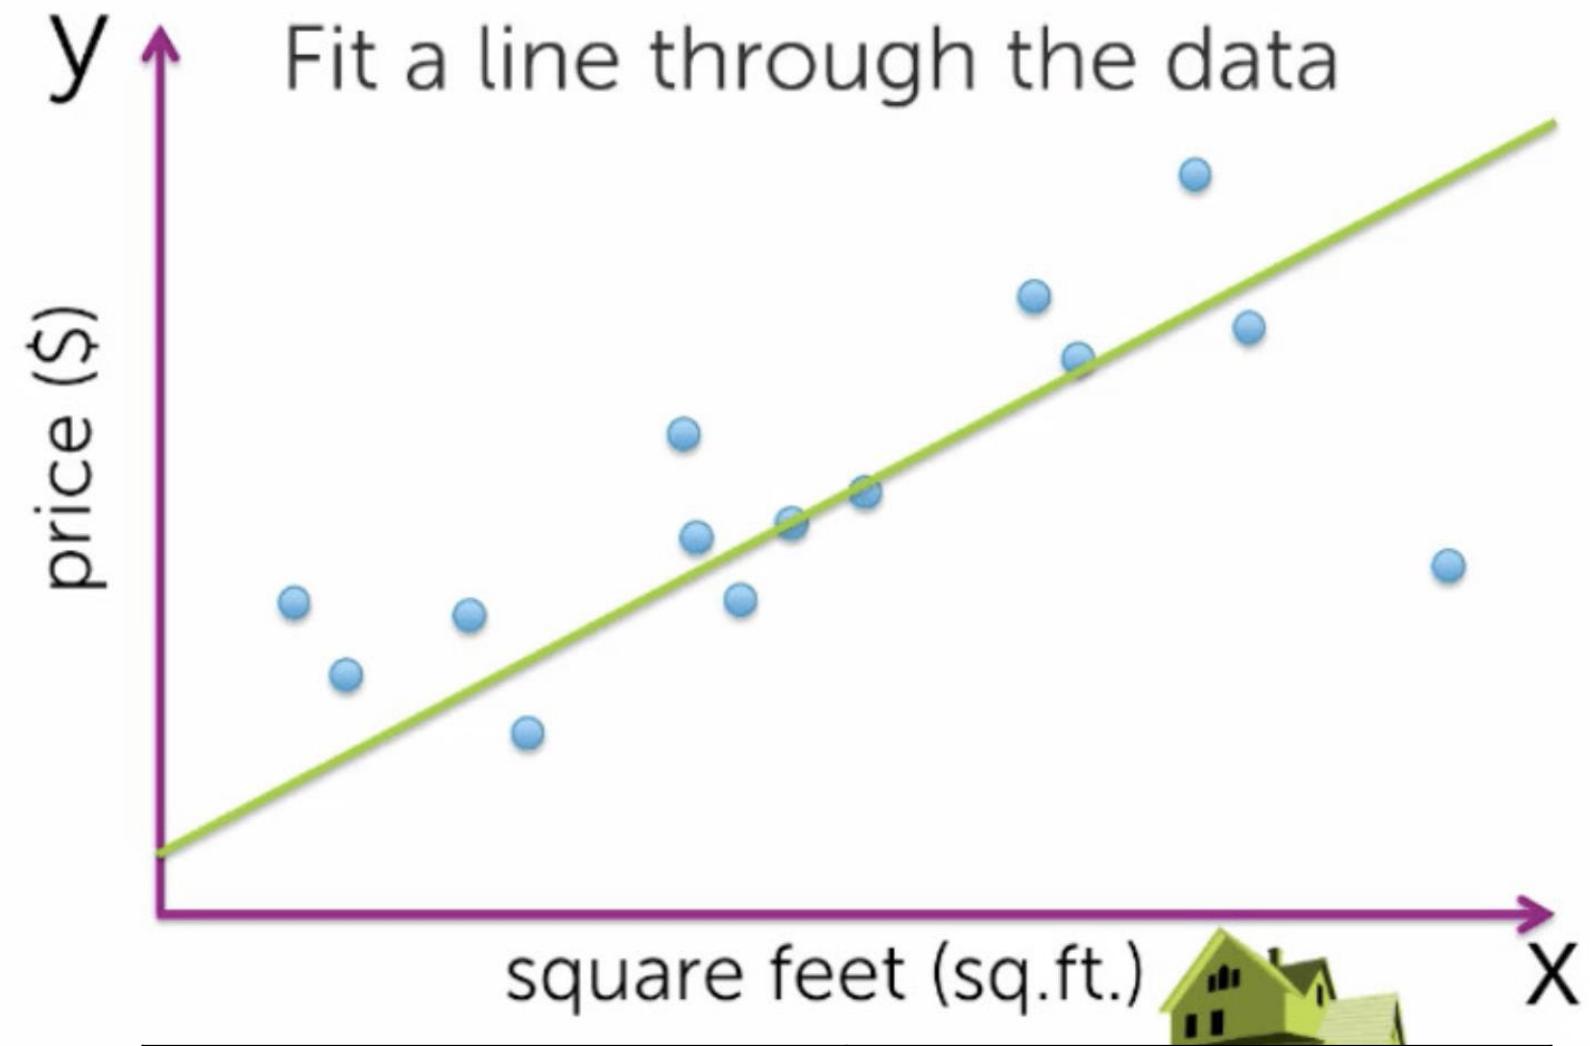
- It learns directly from raw data
- No need for manual feature engineering like in classic ML.
- Examples:
  - Raw pixels → image classifier
  - Raw audio → speech recognizer
  - Raw text → large language model
- It gives state-of-the-art performance
  - Computer vision
  - NLP
  - Speech recognition
  - Recommendation systems
  - Autonomous vehicles
  - Medical imaging
  - Because multiple layers allow **complex patterns** to be captured easily.

# Linearity



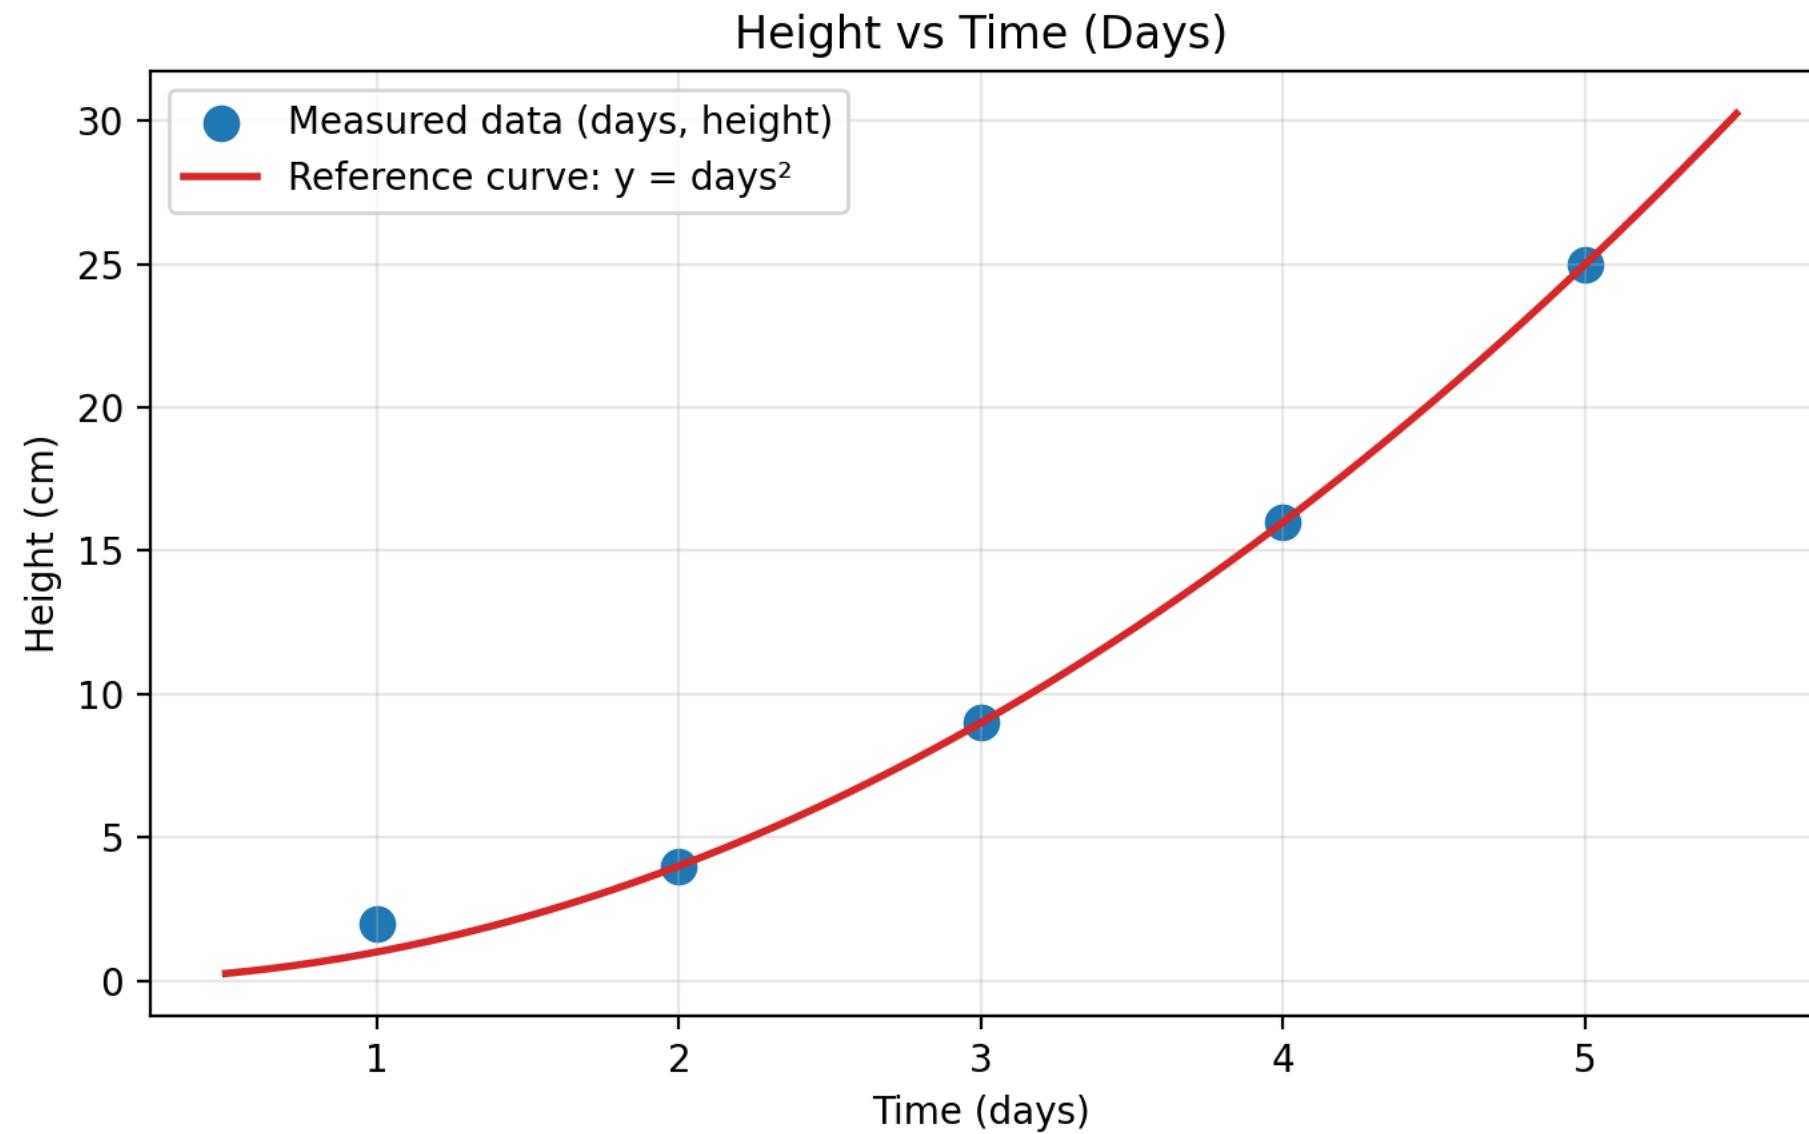
# Linear Data

- When we say data is **linear** or **non-linear**, we're talking about **how the features relate to the target/output**, not about the data points themselves.
- Let me show you with examples.
- **Example 1:** Linear relationship
- **Scenario:** Predicting house price from house size.
- **Pattern:** If you double the size, the price doubles.  
Linear!



# Non Linear Data

- **Scenario:** Predicting plant growth over time.
- **Pattern:** Height  $\approx$  (days) $^2$ .
- Not linear, it's quadratic!
- A quadratic equation is a formula (like  $y = x^2$ ).
- A parabola is the shape you get when you draw (graph) that quadratic equation.
- **Quadratic =** equation
- **Parabola =** graph (curve)
- **Analogy:** Song (equation)  $\rightarrow$  Sound (graph)
- You sing a song  $\rightarrow$  it produces sound
- You write a quadratic  $\rightarrow$  it produces a parabola



Time (days)	Height (cm)
1	2
2	4
3	9
4	16
5	25

# Real-world intuition

---

Linear relationships (common in simple scenarios)

- **Distance** = speed × time.
- **Total cost** = unit price × quantity.
- **Simple physics:** Force = mass × acceleration.

Non-linear relationships (very common in real problems)

- **Biology:** Population growth (exponential), enzyme kinetics (saturation curves).
  - **Finance:** Compound interest ( $A = P(1 + r)^t$ ), option pricing.
  - **Vision:** A pixel's contribution to "is this a cat?" depends on its neighbors and context (highly non-linear).
  - **Language:** Word meaning in context (depends on surrounding words in complex ways).
- 



---

# If Classical ML can Handle Nonlinear Patterns, then why do we Need Deep Learning?



# ML also handles non-linearity

---

- Traditional ML can model non-linear relationships using different mechanisms:

A) Trees / Ensembles (Random Forest, Gradient Boosting)

- Learn **if/then splits** and interactions automatically.
- Very strong on **tabular/structured data**.
- Often win on business datasets with columns like age, income, category, etc.

B) Kernel methods Support Vector Machines (SVM) with a Radial Basis Function (RBF)

- Create non-linear decision boundaries via kernels.

- Good for some medium-sized datasets.

C) kNN

- Non-parametric; predicts from nearest examples.
- Can model complex boundaries but scales poorly.

**Neural networks** also model non-linearity, but the key difference is **how they learn it and where they shine**.



# Main differences (practical + conceptual)

---

Feature engineering vs representation learning • Example:

- **Traditional ML (often):**
  - Images: pixels → edges → shapes → objects
  - Text: tokens → syntax/semantics → meaning
- Works great if you give it good features.
- You frequently do manual feature work:
  - transformations (log, square)
  - aggregations
  - domain-crafted signals
- **Neural networks / Deep learning:**
  - Do **representation learning**: learn features automatically from raw inputs.



# Best Data Types (Where each Dominates)

---

Data type	Traditional ML (trees/boosting)	Deep learning
Tabular structured (columns)	Often best / strongest baseline	Sometimes good, often not best unless huge data
Images	Weak without hand-crafted features	Usually best (CNNs/ViTs)
Text / NLP	Weak with bag-of-words vs modern tasks	Usually best (Transformers)
Audio / speech	Limited	Usually best
Video	Limited	Usually best

# A simple way to say it

---

- **Machine learning** = all learning models (linear, trees, SVM, neural nets, etc.).
- **Neural networks** = one ML approach that builds non-linear functions by stacking layers.
- **Deep learning** = neural networks with many layers that can learn features from raw data.



# Manual Features Vs Automatic Features

---

## Traditional ML pipeline (manual feature engineering)

- **Step 1 Input:** Raw image (pixels), e.g.  $128 \times 128 \times 3$ .
- **Step 2 Manually design features** (human-chosen)
  - You don't feed raw pixels directly (or if you do, it performs poorly). You first extract engineered features such as:
- **HOG (Histogram of Oriented Gradients):** Captures edge directions and local shapes (good for object outlines).
- **SIFT / SURF / ORB keypoints:** Detects distinctive points and describes local patches.
- **Color histograms:** How much of each color appears (sometimes useful, sometimes misleading).
- **Texture features (LBP, Gabor filters):** Captures fur-like texture patterns.

- After this, each image becomes a fixed feature vector:

$$x = \phi(\text{image})$$

- **Step 3:** Train a “classic” model: SVM, Logistic Regression, Random Forest, kNN

- **What's manual here?**

You chose the feature extractor  $\phi \left( \frac{\text{HOG}}{\text{SIFT}} \dots \right)$  and its settings.

- **Limitation:**

If the hand-designed features don't capture what matters (pose, lighting, background, breed variety), performance is limited.



# Automatic Feature Learning

---

- **Step 1 Input:** Raw pixels (possibly resized/normalized).
- **Step 2 Neural network learns features automatically:** Use a CNN (e.g., ResNet) or Vision Transformer.
- The network learns a hierarchy:
- **Early layers:** edges, corners (similar to what HOG tries to capture)
- **Middle layers:** textures, patterns (fur, whiskers), simple shapes
- **Deeper layers:** parts (eyes, ears, snout)
- **Final layers:** whole object concept (cat vs dog)
- So the model learns:  
 $\phi(\cdot)$  *automatically during training*
- **Step 3 Classifier head:** Final layer predicts cat/dog.
- **What's “manual” here?**  
Mostly basic preprocessing:
  - resize/crop
  - normalize pixel values
  - maybe data augmentation (flip, rotate)
  - But you did **not** manually define edges/textures/keypoints, those are learned.

# Learning vs Inference: Conceptual Distinction

---

- **Learning (Training phase):**
  - Model parameters are adjusted to fit training data
  - Uses optimization algorithms (e.g., gradient descent)
  - Computationally expensive, off-line (often on GPUs)
- **Inference (Prediction phase):**
  - Using a trained model to make predictions on **new** inputs
  - Usually faster, runs on servers or edge devices
- **Mental model:**
  - Learning = *studying for the exam*
  - Inference = *answering questions during the exam*



# Generalization: The Central Goal

---

- **Generalization:**
  - A model's ability to perform well on **unseen data** drawn from the same distribution as the training data.
- **Good learning:**
  - Not just memorizing training data
  - Captures underlying patterns/structure
- **Evaluation:**
  - Train/validation/test splits
  - Metrics (accuracy, error, loss)

# Overfitting and Underfitting

- **Underfitting:**

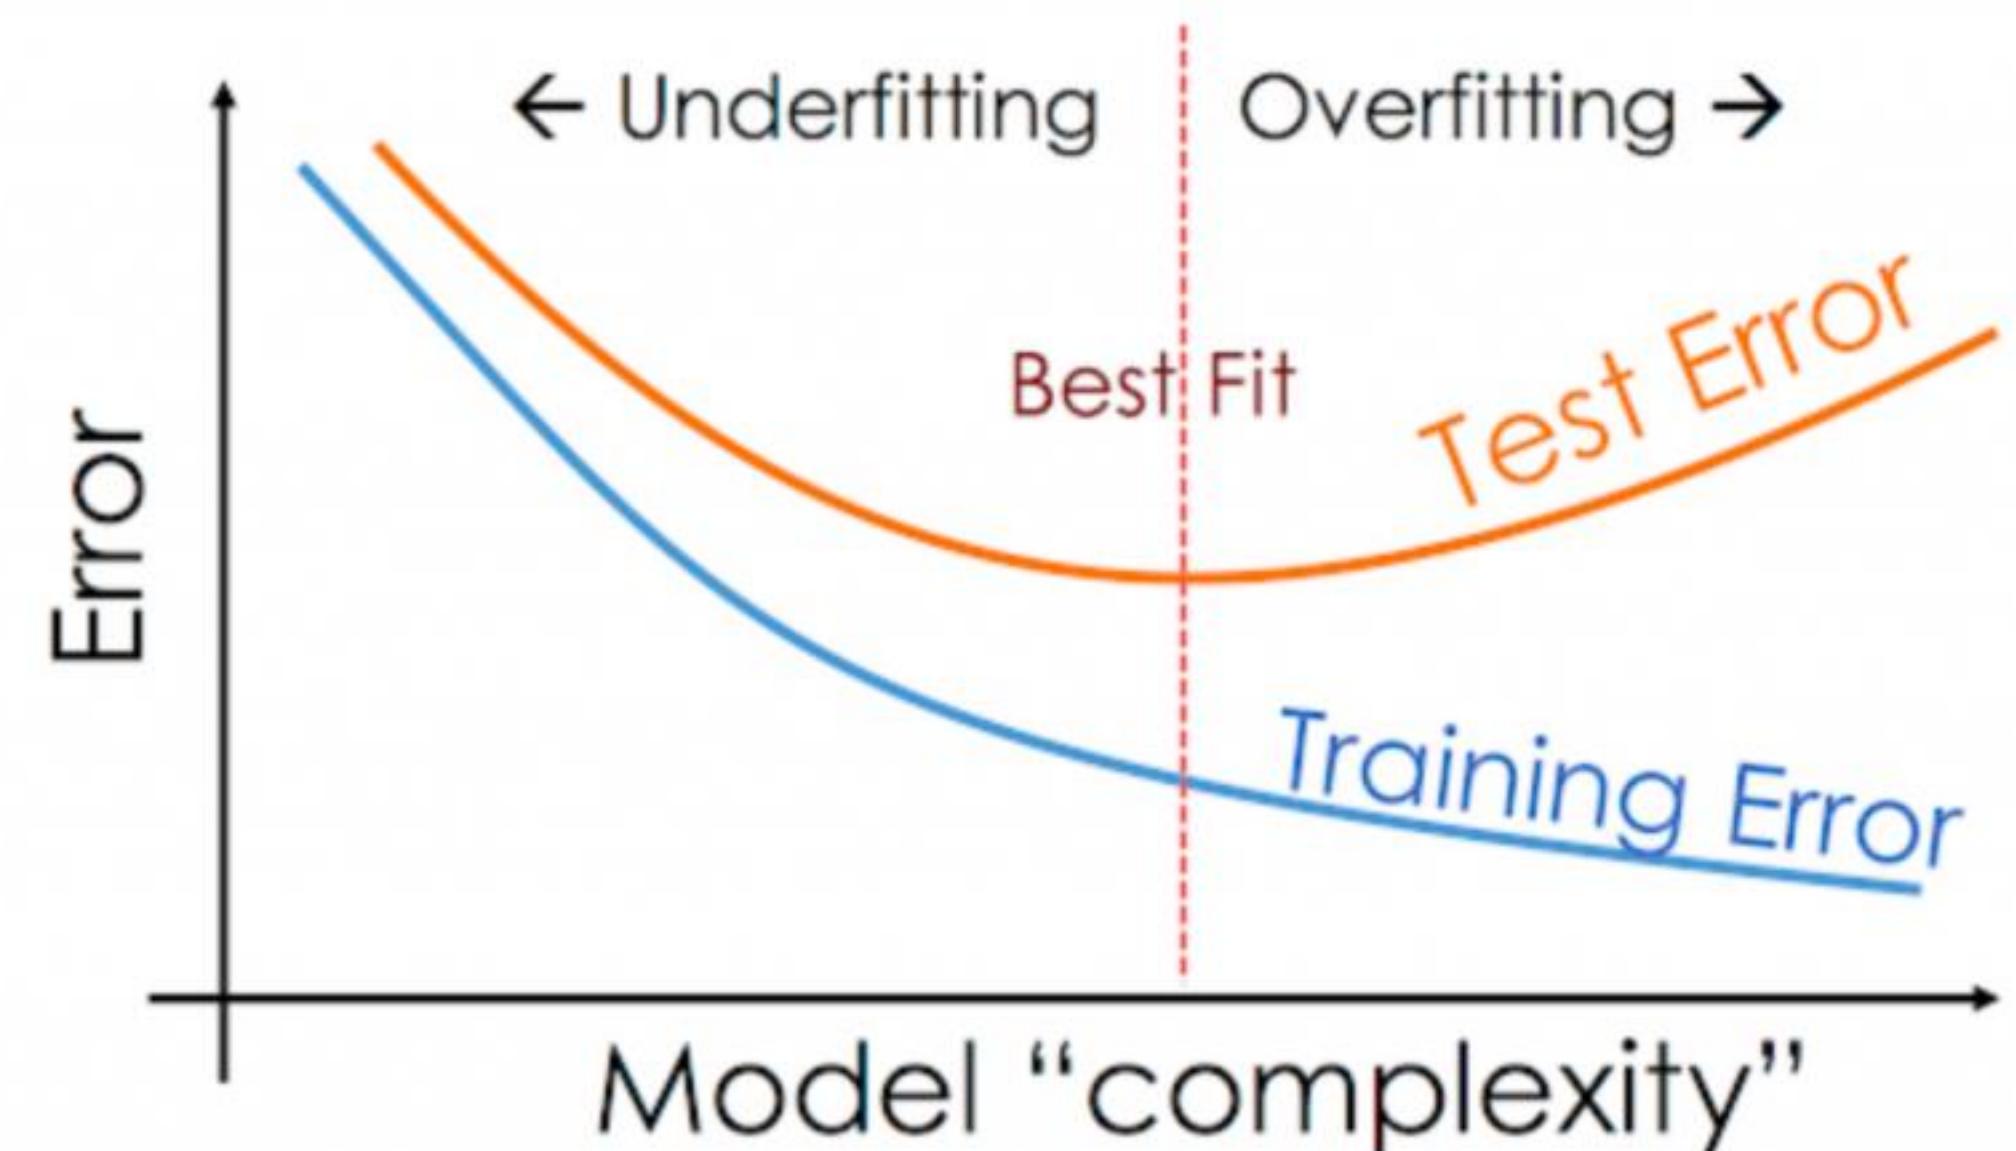
- Model is too simple to capture the data patterns
- High error on both training and test data

- **Overfitting:**

- Model fits noise in training data
- Low training error; high test error

- **Just right (good fit):**

- Low training error and low test error



# Controlling Overfitting

---

- Techniques to improve generalization:
- More data / better data
- Regularization (e.g., weight decay, dropout)
- Early stopping based on validation performance
- Simpler architectures / fewer parameters
- Data augmentation (esp. in vision)
- Important message:
- Deep networks are powerful; **capacity** must be controlled carefully.



# Types of Learning Paradigms

---

- Fully supervised learning
- Unsupervised learning
- Semi-supervised learning
- Self-supervised learning
- Reinforcement learning
- **Axis of difference:**
- Type and amount of **supervision / feedback** available.

# Fully Supervised Learning

---

- **Setting:**
- You have input-output pairs:  $\{(x_i, y_i)\}$
- Goal: learn a function  $f$  that maps  $x$  to  $y$ .
- **Examples:**
  - Image classification (image  $\rightarrow$  label)
  - Sentiment analysis (text  $\rightarrow$  sentiment)
  - Speech recognition (audio  $\rightarrow$  transcript)
- **Pros & cons:**
  - **Pros:** Strong performance when labels are abundant
  - **Cons:** Labelling is expensive/time-consuming



# Regression vs Classification (Supervised)

- **Regression:** Predict **continuous** values
- House price prediction
- Stock price forecasting
- Temperature prediction
- **Common loss:** Mean Squared Error (MSE)



## Regression

What is the temperature going to be tomorrow?

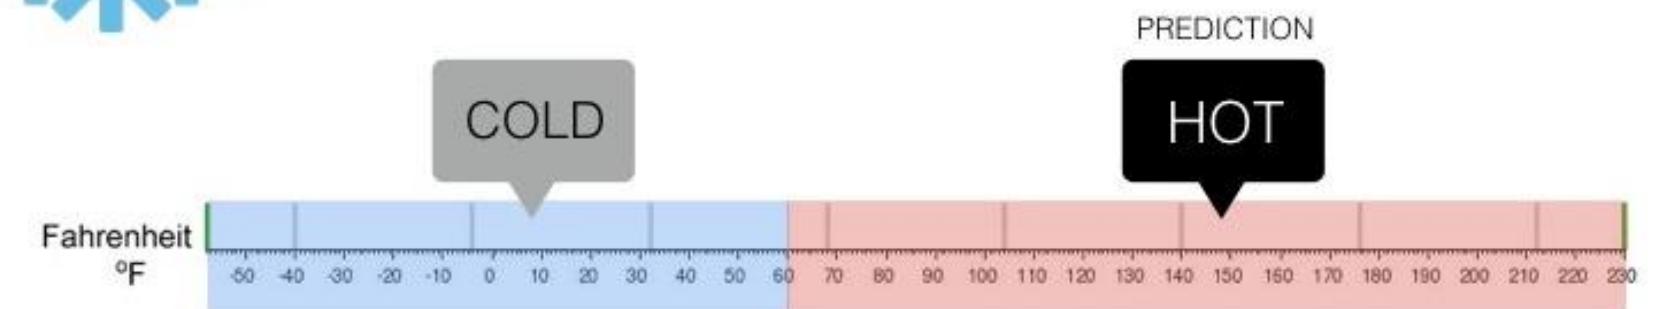


- **Classification:** Predict **discrete** labels
- Cat vs dog
- Spam vs not spam
- Multi-class object categories
- **Common loss:** Cross-entropy



## Classification

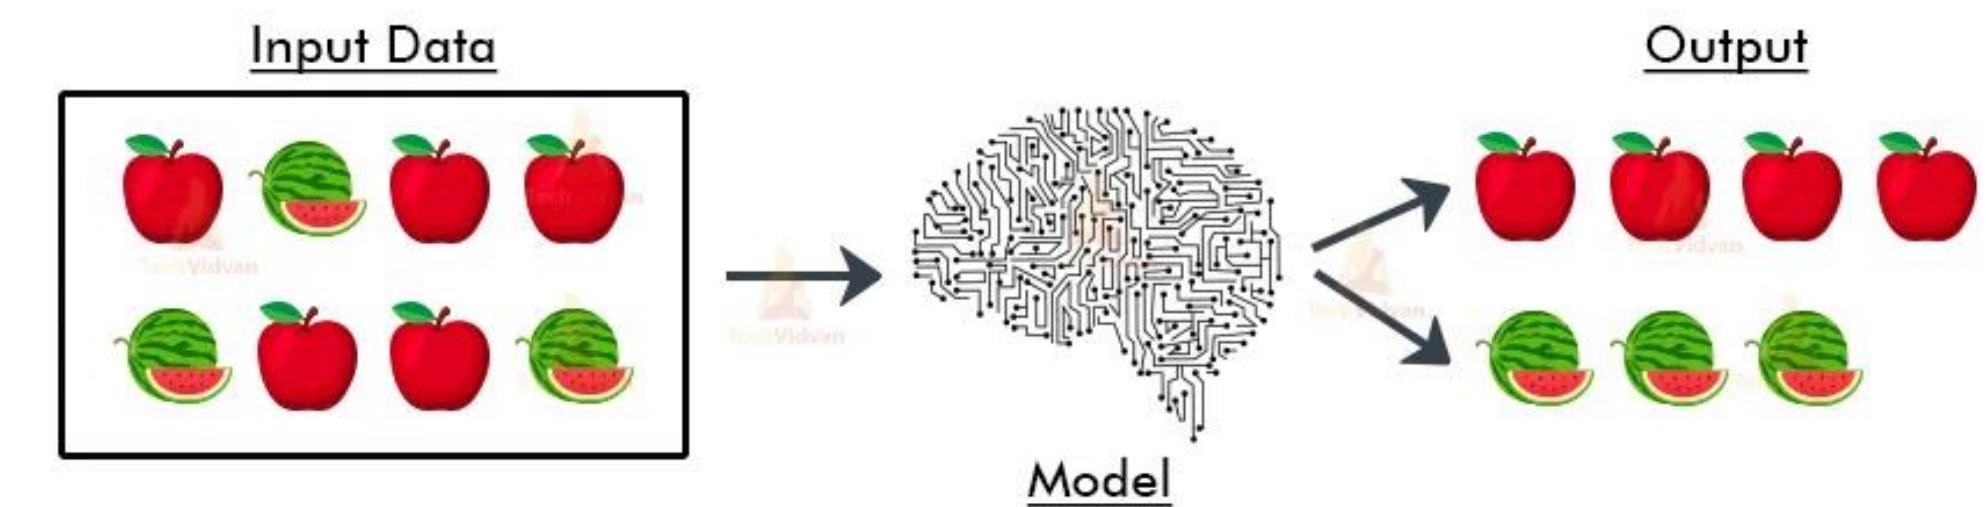
Will it be Cold or Hot tomorrow?



# Unsupervised Learning

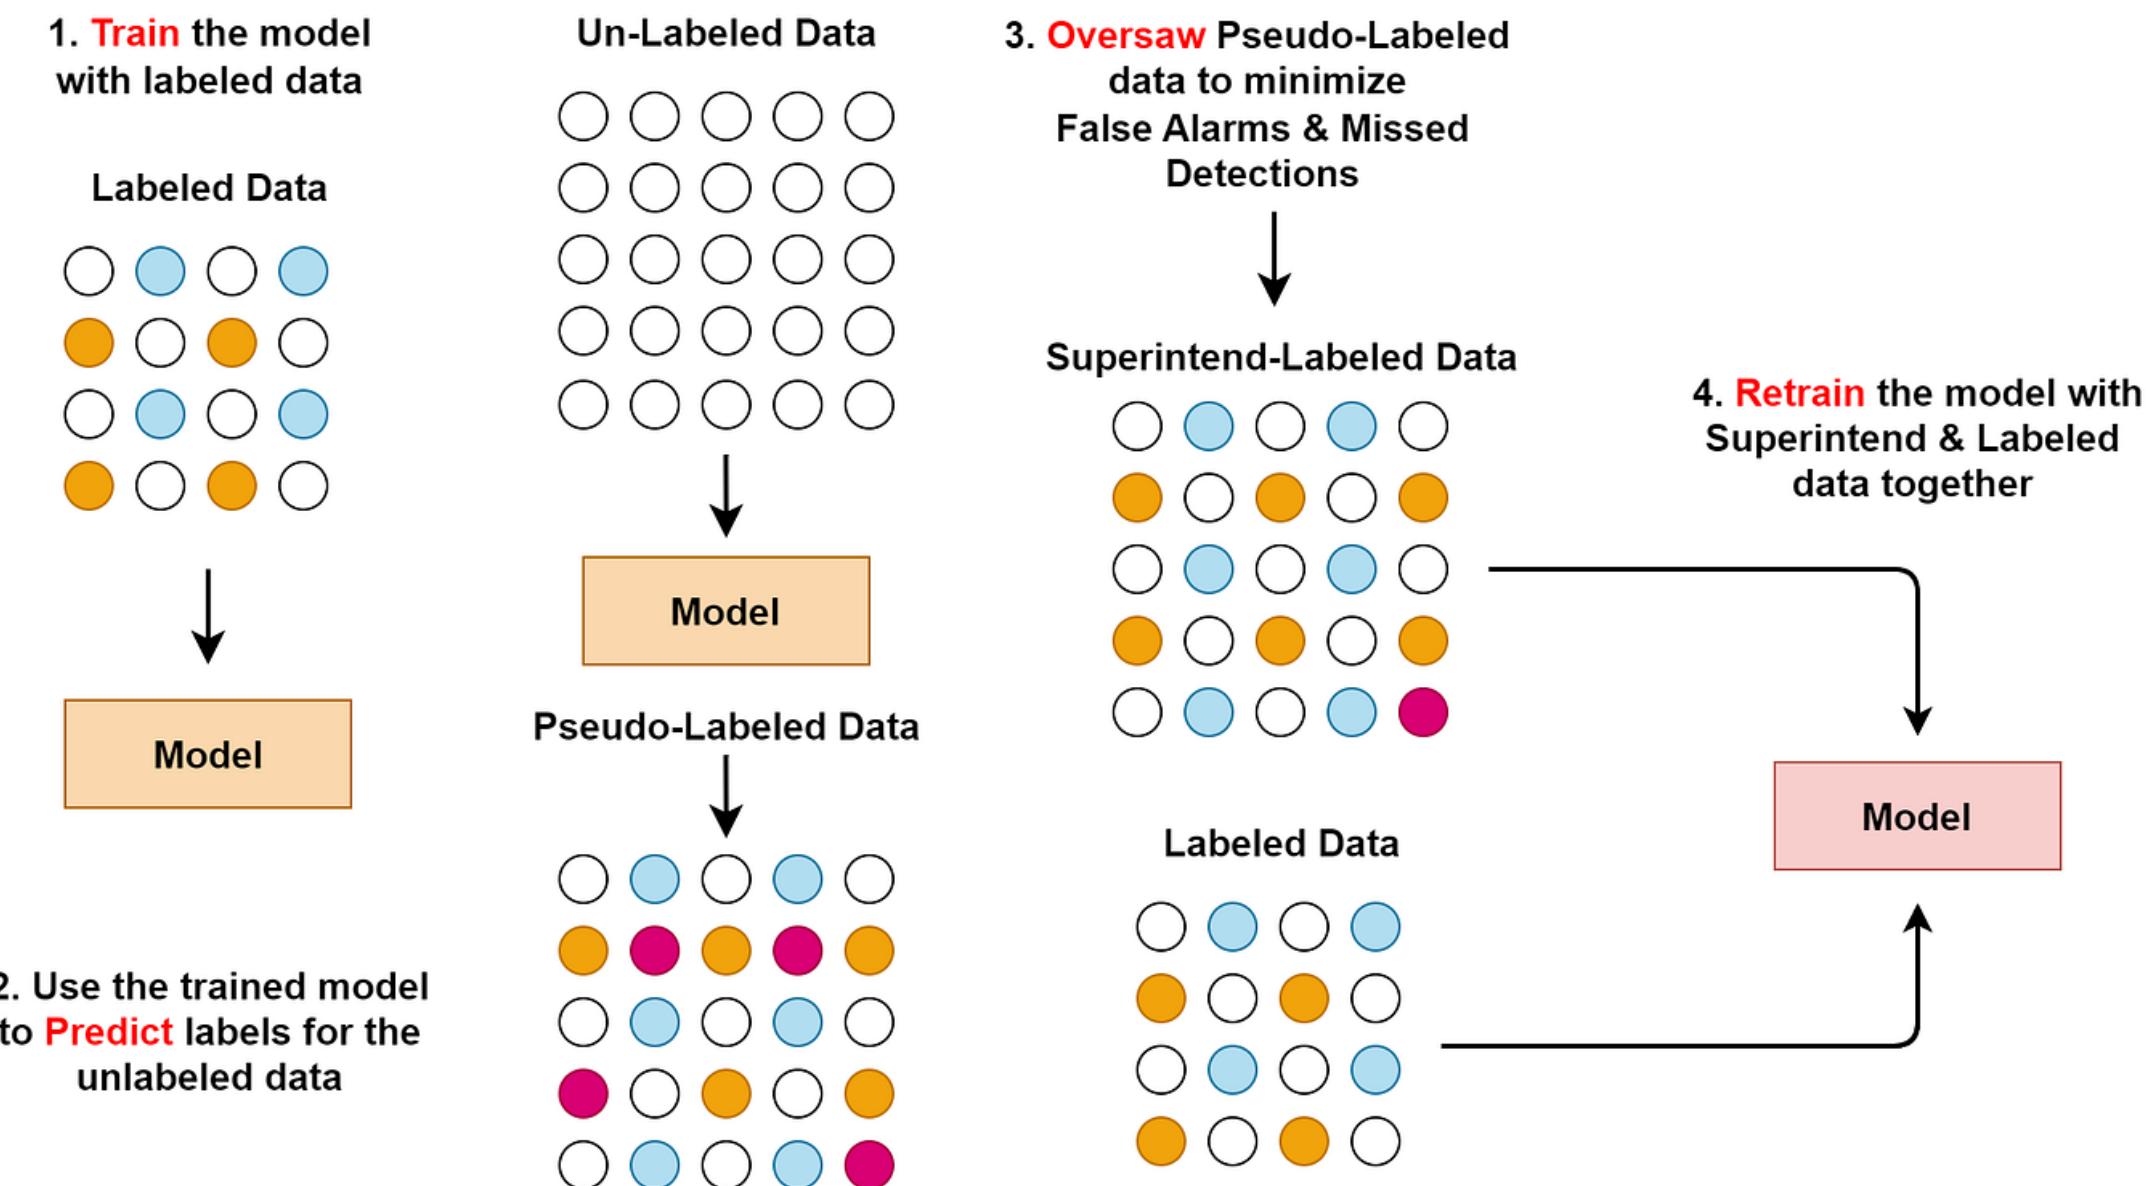
---

- Only input data  $\{x_i\}$ , **no labels**.
- Discover structure in data:
- Clustering (group similar examples)
- Dimensionality reduction (compress data)
- Density estimation, anomaly detection
- **Examples:**
- k-means clustering, PCA, autoencoders
- **Motivation:**
- Labels are scarce, but unlabeled data is cheap and abundant.



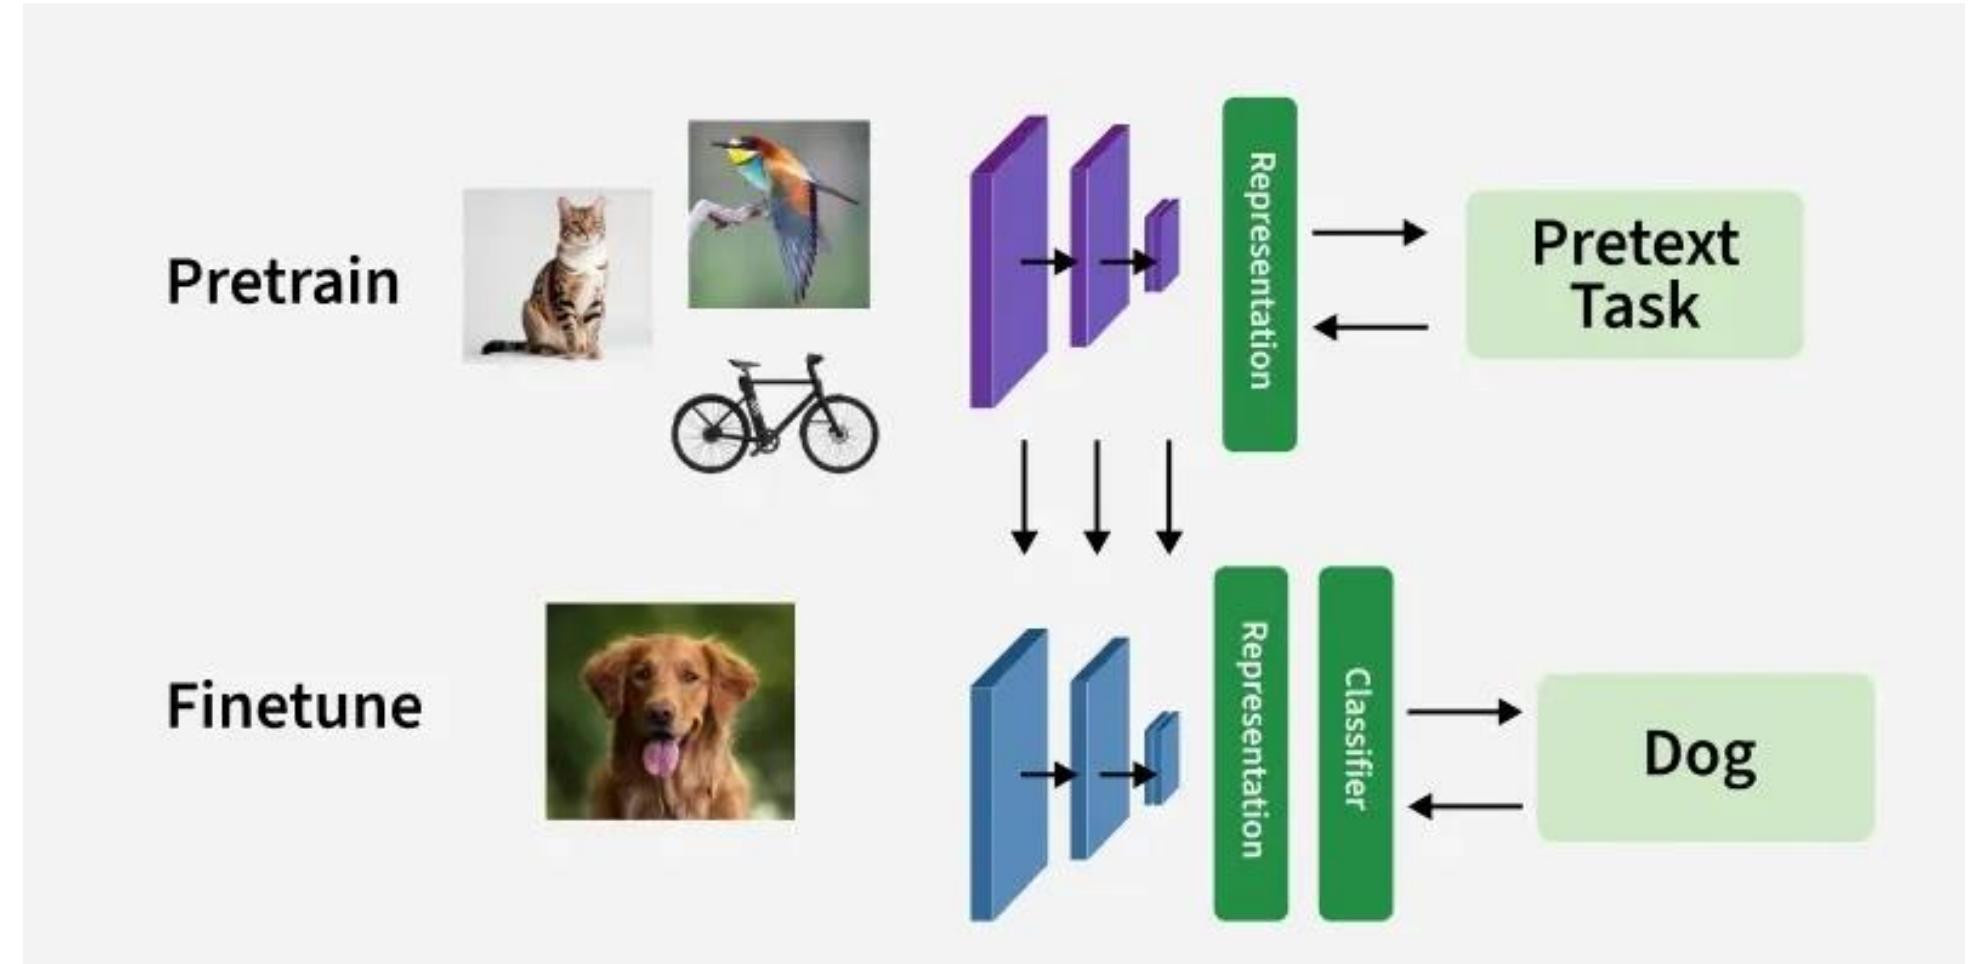
# Semi-Supervised Learning

- Small amount of labeled data + large amount of unlabeled data
- Leverage unlabeled data to improve performance vs purely supervised learning with few labels
- **Approaches:**
  - Consistency regularization
  - Pseudo-labeling
  - Graph-based methods
- **Use cases:**
  - Medical imaging, low-resource languages, any domain where labeling is costly.



# Self-Supervised Learning

- Create **pseudo-labels** from the data itself; solve pretext tasks.
- **Examples of pretext tasks:**
  - Predict missing part of an image or sentence
  - Next-token prediction in language
  - Contrastive learning: distinguish true pairs vs mismatched pairs
- **Why important:**
  - Enables learning powerful representations without manual labels
  - Foundation of many modern **large language models** and vision models.



# Reinforcement Learning: Delayed Feedback

---

- **Delayed Feedback:** You do many actions, but you get the reward **later**, not immediately.  
No instant feedback.
- **Supervised Learning:**
  - You give input → model predicts → you immediately give correct label.  
Instant feedback.
- **Reinforcement Learning:**
  - Agent takes **many steps**, then gets reward at the end.
  - **Example**
    - Agent walks 10 steps to reach the goal.
    - It only gets the +10 reward **after reaching the goal**,  
NOT after each step.
    - **This is delayed feedback.**

# Reinforcement Learning : Sparse Feedback

---

- Rewards are rare.  
Most steps give **no reward**.
- **Supervised Learning:**
- Every example has a label.
- Feedback for every data point.
- **Reinforcement Learning:**
- Most steps = no reward.  
Only sometimes do you get reward.
- Example:
  - Walk → 0 reward
  - Walk → 0 reward
  - Walk → 0 reward
  - Reach goal → +10 reward
- **This is sparse feedback.**

# Reinforcement Learning : Non-IID Data

---

- IID = Independent and Identically Distributed
- The data in RL is **dependent** on previous actions.

You create your own data by moving in the environment.

- **Supervised Learning:**
- Each training example is separate and independent.

Example:

- “I love this movie”
- “This is bad”

- “Amazing!”  
They don’t depend on each other.
- Data is IID.
- **Reinforcement Learning:**
- State 1 leads to State 2, which leads to State 3...  
Each next step depends on the previous step.
- Example: If agent moves right → next state is determined by that action.
- **Data is NOT independent.**  
**This is non-IID.**

# What is Intelligence? Perspectives

---

- Defining Intelligence (Difficult Problem)
- **No single universally accepted definition.**
- **Informal views:**
  - Ability to learn, reason, adapt, and solve novel problems
  - Ability to achieve goals across a wide range of environments
- **Human vs Machine intelligence:**
  - Humans: embodied, emotional, social, grounded in physical world
  - Machines: currently narrow, task-specific, data-driven.

# Perspectives on Intelligence

---

- **Psychological Perspective:**
    - Cognitive abilities: memory, reasoning, problem-solving, creativity
    - Measured historically via IQ tests (with many critiques)
  - **AI perspective:**
    - Performance on benchmarks/tasks
    - Rational agents maximizing expected utility
  - **Philosophical perspective:**
    - Consciousness, understanding, intentionality
-

# Is Deep Learning “Intelligent”?

---

- **Strengths:**
- High performance on specific tasks
- Can approximate complex functions, recognize patterns beyond human capability
- **Weaknesses:**
- Lack of robust common sense and causal reasoning
- Brittle outside training distribution
- Limited interpretability
- Deep learning systems exhibit *narrow* forms of intelligence; broader intelligence remains open research.

---

# Real-World Applications Across Domains



# Computer Vision

---

- **Tasks:**

- Image classification, object detection, segmentation
- Face recognition, activity recognition

- **Applications:**

- Self-driving cars (perception stack)
- Medical image analysis (tumor detection)
- Surveillance and security
- Industrial inspection and quality control



# Natural Language Processing (NLP)

---

- **Tasks:**

- Machine translation, summarization
- Question answering, chatbots
- Sentiment analysis, information extraction

- **Applications:**

- Virtual assistants and customer support
- Legal and financial document analysis
- Search engines and recommendation systems



# Speech, Audio & Multimodal Applications

---

- **Speech & audio:**

- Speech recognition, text-to-speech
- Speaker identification, emotion recognition
- Music recommendation, sound event detection

- **Multimodal AI:**

- Combining text, image, audio, video
- Examples: video captioning, visual Q&A, AR/VR assistants



# Robotics and Control

---

- **Applications:**

- Industrial robots and cobots
- Drones and autonomous vehicles
- Household robots (vacuum cleaners, assistive devices)

- **Role of deep learning:**

- Perception (seeing the world)
- Policy learning (RL) for complex control tasks
- Sensor fusion and decision-making.



# Finance, Business & Recommendation

---

- **Finance:**
  - Algorithmic trading, risk prediction, fraud detection
- **Business:**
  - Demand forecasting, inventory optimization
  - Customer segmentation, churn prediction
- **Recommender systems:**
  - E-commerce product recommendations
  - Personalized content feeds
- **Value proposition:**
  - Data-driven decisions, improved efficiency, personalization.



# Science & Medicine

---

- **Scientific discovery:**

- Protein structure prediction
- Materials discovery; drug design
- Physics simulations and surrogates

- **Medicine:**

- Predictive models for disease risk
- Personalized treatment recommendations
- Clinical decision support systems



# Paper Writing

---

- A typical paper must consist of following main section
  - Abstract
    - a. Should be one paragraph explaining what is being presented in the paper
  - Introduction
    - a. A short explanation of the problem and reviews of the existing research.
    - b. You must download and review at least 10-15 papers (Journal Papers are highly encouraged)
  - Data Set
    - a. Which data has been used for this study
  - Methodology
    - a. What method(s) has been implemented
  - Experimental Results
    - a. Explanation of the results obtained. Should include figures, tables etc.
  - Conclusion
    - a. Summarize what you did overall
  - References
    - a. Bibliography of which papers were reviewed or used in the paper



# Getting Started with Research

---

- Step 1. Select a Topic
  - Your interests and Skills
  - Supervisor / Project
  - Availability of resources
- Step 2. Search for Relevant Information
  - IEEE Explore
  - Springer
  - Scopus
  - Web of Science
  - Google Scholar
- Step 3. Organize Resources
  - Zotero
  - EndNote
  - Mendeley
  - JabRef
  - License
  - MarginNote
- Step 4. Write Paper
  - Outline the Paper
  - Take Notes
  - Avoid Plagiarism



---

# Your Tasks

**I will review this in the same day next class and assign the corresponding marks.**

